

# class 08 mini project

Lawrence Adhinatha

2023-02-02

In today's mini-project we will explore a complete analysis using the unsupervised learning techniques covered in class (clustering and PCA for now).

The data itself comes from the Wisconsin Breast Cancer Diagnostic data set FNA biopsy data.

## Exploratory Data Analysis

### Data import

```
wisc.df <- read.csv("WisconsinCancer.csv", row.names=1)
head(wisc.df)
```

```
##      diagnosis radius_mean texture_mean perimeter_mean area_mean
## 842302         M      17.99       10.38         122.80      1001.0
## 842517         M      20.57       17.77         132.90      1326.0
## 84300903        M      19.69       21.25         130.00      1203.0
## 84348301         M      11.42       20.38          77.58       386.1
## 84358402         M      20.29       14.34         135.10      1297.0
## 843786          M      12.45       15.70          82.57       477.1
##      smoothness_mean compactness_mean concavity_mean concave.points_mean
## 842302          0.11840         0.27760         0.3001          0.14710
## 842517          0.08474         0.07864         0.0869          0.07017
## 84300903         0.10960         0.15990         0.1974          0.12790
## 84348301         0.14250         0.28390         0.2414          0.10520
## 84358402         0.10030         0.13280         0.1980          0.10430
## 843786          0.12780         0.17000         0.1578          0.08089
##      symmetry_mean fractal_dimension_mean radius_se texture_se perimeter_se
## 842302          0.2419         0.07871      1.0950      0.9053          8.589
## 842517          0.1812         0.05667      0.5435      0.7339          3.398
## 84300903         0.2069         0.05999      0.7456      0.7869          4.585
## 84348301         0.2597         0.09744      0.4956      1.1560          3.445
## 84358402         0.1809         0.05883      0.7572      0.7813          5.438
## 843786          0.2087         0.07613      0.3345      0.8902          2.217
##      area_se smoothness_se compactness_se concavity_se concave.points_se
## 842302     153.40      0.006399      0.04904      0.05373      0.01587
## 842517      74.08      0.005225      0.01308      0.01860      0.01340
## 84300903     94.03      0.006150      0.04006      0.03832      0.02058
## 84348301     27.23      0.009110      0.07458      0.05661      0.01867
## 84358402     94.44      0.011490      0.02461      0.05688      0.01885
```

```
## 843786      27.19      0.007510      0.03345      0.03672      0.01137
##      symmetry_se fractal_dimension_se radius_worst texture_worst
## 842302      0.03003      0.006193      25.38      17.33
## 842517      0.01389      0.003532      24.99      23.41
## 84300903      0.02250      0.004571      23.57      25.53
## 84348301      0.05963      0.009208      14.91      26.50
## 84358402      0.01756      0.005115      22.54      16.67
## 843786      0.02165      0.005082      15.47      23.75
##      perimeter_worst area_worst smoothness_worst compactness_worst
## 842302      184.60      2019.0      0.1622      0.6656
## 842517      158.80      1956.0      0.1238      0.1866
## 84300903      152.50      1709.0      0.1444      0.4245
## 84348301      98.87      567.7      0.2098      0.8663
## 84358402      152.20      1575.0      0.1374      0.2050
## 843786      103.40      741.6      0.1791      0.5249
##      concavity_worst concave.points_worst symmetry_worst
## 842302      0.7119      0.2654      0.4601
## 842517      0.2416      0.1860      0.2750
## 84300903      0.4504      0.2430      0.3613
## 84348301      0.6869      0.2575      0.6638
## 84358402      0.4000      0.1625      0.2364
## 843786      0.5355      0.1741      0.3985
##      fractal_dimension_worst
## 842302      0.11890
## 842517      0.08902
## 84300903      0.08758
## 84348301      0.17300
## 84358402      0.07678
## 843786      0.12440
```

Remove the diagnosis column and keep it in a separate vector for later.

```
diagnosis <- as.factor(wisc.df[,1])
wisc.data <- wisc.df[,-1]
head(wisc.data)
```

```
##      radius_mean texture_mean perimeter_mean area_mean smoothness_mean
## 842302      17.99      10.38      122.80      1001.0      0.11840
## 842517      20.57      17.77      132.90      1326.0      0.08474
## 84300903      19.69      21.25      130.00      1203.0      0.10960
## 84348301      11.42      20.38      77.58      386.1      0.14250
## 84358402      20.29      14.34      135.10      1297.0      0.10030
## 843786      12.45      15.70      82.57      477.1      0.12780
##      compactness_mean concavity_mean concave.points_mean symmetry_mean
## 842302      0.27760      0.3001      0.14710      0.2419
## 842517      0.07864      0.0869      0.07017      0.1812
## 84300903      0.15990      0.1974      0.12790      0.2069
## 84348301      0.28390      0.2414      0.10520      0.2597
## 84358402      0.13280      0.1980      0.10430      0.1809
## 843786      0.17000      0.1578      0.08089      0.2087
##      fractal_dimension_mean radius_se texture_se perimeter_se area_se
## 842302      0.07871      1.0950      0.9053      8.589      153.40
## 842517      0.05667      0.5435      0.7339      3.398      74.08
```

```
## 84300903      0.05999      0.7456      0.7869      4.585      94.03
## 84348301      0.09744      0.4956      1.1560      3.445      27.23
## 84358402      0.05883      0.7572      0.7813      5.438      94.44
## 843786        0.07613      0.3345      0.8902      2.217      27.19
##      smoothness_se compactness_se concavity_se concave.points_se
## 842302      0.006399      0.04904      0.05373      0.01587
## 842517      0.005225      0.01308      0.01860      0.01340
## 84300903      0.006150      0.04006      0.03832      0.02058
## 84348301      0.009110      0.07458      0.05661      0.01867
## 84358402      0.011490      0.02461      0.05688      0.01885
## 843786        0.007510      0.03345      0.03672      0.01137
##      symmetry_se fractal_dimension_se radius_worst texture_worst
## 842302      0.03003      0.006193      25.38      17.33
## 842517      0.01389      0.003532      24.99      23.41
## 84300903      0.02250      0.004571      23.57      25.53
## 84348301      0.05963      0.009208      14.91      26.50
## 84358402      0.01756      0.005115      22.54      16.67
## 843786        0.02165      0.005082      15.47      23.75
##      perimeter_worst area_worst smoothness_worst compactness_worst
## 842302      184.60      2019.0      0.1622      0.6656
## 842517      158.80      1956.0      0.1238      0.1866
## 84300903      152.50      1709.0      0.1444      0.4245
## 84348301      98.87      567.7      0.2098      0.8663
## 84358402      152.20      1575.0      0.1374      0.2050
## 843786        103.40      741.6      0.1791      0.5249
##      concavity_worst concave.points_worst symmetry_worst
## 842302      0.7119      0.2654      0.4601
## 842517      0.2416      0.1860      0.2750
## 84300903      0.4504      0.2430      0.3613
## 84348301      0.6869      0.2575      0.6638
## 84358402      0.4000      0.1625      0.2364
## 843786        0.5355      0.1741      0.3985
##      fractal_dimension_worst
## 842302      0.11890
## 842517      0.08902
## 84300903      0.08758
## 84348301      0.17300
## 84358402      0.07678
## 843786        0.12440
```

Q1. How many observations are in this dataset?

```
nrow(wisc.data)
```

```
## [1] 569
```

There are 569 cells observed in this dataset.

Q2. How many of the observations have a malignant diagnosis?

```
table(diagnosis)
```

```
## diagnosis
##   B   M
## 357 212
```

212 cells have a malignant diagnosis.

Q3. How many variables/features in the data are suffixed with `_mean`?

First find the column names

```
colnames(wisc.data)
```

```
## [1] "radius_mean"      "texture_mean"
## [3] "perimeter_mean"   "area_mean"
## [5] "smoothness_mean"  "compactness_mean"
## [7] "concavity_mean"    "concave.points_mean"
## [9] "symmetry_mean"     "fractal_dimension_mean"
## [11] "radius_se"         "texture_se"
## [13] "perimeter_se"      "area_se"
## [15] "smoothness_se"     "compactness_se"
## [17] "concavity_se"      "concave.points_se"
## [19] "symmetry_se"       "fractal_dimension_se"
## [21] "radius_worst"      "texture_worst"
## [23] "perimeter_worst"   "area_worst"
## [25] "smoothness_worst"  "compactness_worst"
## [27] "concavity_worst"   "concave.points_worst"
## [29] "symmetry_worst"    "fractal_dimension_worst"
```

Next, I need to search within the column names for “`_mean`” pattern. The `grep()` function might be of use.

```
length(grep("_mean", colnames(wisc.data)))
```

```
## [1] 10
```

There are 10 features in the data suffixed with `_mean`.

Q. How many dimensions are in this dataset?

```
ncol(wisc.data)
```

```
## [1] 30
```

## Principal Component Analysis

First, check if the variables are all on the same scale: Do we need to scale the data before performing PCA?

```
round(apply(wisc.data, 2, sd), 3)
```

```
##           radius_mean      texture_mean      perimeter_mean
##           3.524          4.301          24.299
##           area_mean      smoothness_mean      compactness_mean
##           351.914        0.014          0.053
##           concavity_mean  concave.points_mean      symmetry_mean
##           0.080          0.039          0.027
## fractal_dimension_mean      radius_se      texture_se
##           0.007          0.277          0.552
##           perimeter_se      area_se      smoothness_se
##           2.022          45.491          0.003
##           compactness_se      concavity_se      concave.points_se
##           0.018          0.030          0.006
##           symmetry_se      fractal_dimension_se      radius_worst
##           0.008          0.003          4.833
##           texture_worst      perimeter_worst      area_worst
##           6.146          33.603          569.357
##           smoothness_worst      compactness_worst      concavity_worst
##           0.023          0.157          0.209
##           concave.points_worst      symmetry_worst      fractal_dimension_worst
##           0.066          0.062          0.018
```

Looks like we need to scale.

```
# Perform PCA on wisc.data
wisc.pr <- prcomp(wisc.data, scale=TRUE)
summary(wisc.pr)
```

```
## Importance of components:
##           PC1      PC2      PC3      PC4      PC5      PC6      PC7
## Standard deviation  3.6444 2.3857 1.67867 1.40735 1.28403 1.09880 0.82172
## Proportion of Variance 0.4427 0.1897 0.09393 0.06602 0.05496 0.04025 0.02251
## Cumulative Proportion 0.4427 0.6324 0.72636 0.79239 0.84734 0.88759 0.91010
##           PC8      PC9      PC10      PC11      PC12      PC13      PC14
## Standard deviation  0.69037 0.6457 0.59219 0.5421 0.51104 0.49128 0.39624
## Proportion of Variance 0.01589 0.0139 0.01169 0.0098 0.00871 0.00805 0.00523
## Cumulative Proportion 0.92598 0.9399 0.95157 0.9614 0.97007 0.97812 0.98335
##           PC15      PC16      PC17      PC18      PC19      PC20      PC21
## Standard deviation  0.30681 0.28260 0.24372 0.22939 0.22244 0.17652 0.1731
## Proportion of Variance 0.00314 0.00266 0.00198 0.00175 0.00165 0.00104 0.0010
## Cumulative Proportion 0.98649 0.98915 0.99113 0.99288 0.99453 0.99557 0.9966
##           PC22      PC23      PC24      PC25      PC26      PC27      PC28
## Standard deviation  0.16565 0.15602 0.1344 0.12442 0.09043 0.08307 0.03987
## Proportion of Variance 0.00091 0.00081 0.0006 0.00052 0.00027 0.00023 0.00005
## Cumulative Proportion 0.99749 0.99830 0.9989 0.99942 0.99969 0.99992 0.99997
##           PC29      PC30
## Standard deviation  0.02736 0.01153
## Proportion of Variance 0.00002 0.00000
## Cumulative Proportion 1.00000 1.00000
```

Q4. From your results, what proportion of the original variance is captured by the first principal components (PC1)?

44.27%

Q5. How many principal components (PCs) are required to describe at least 70% of the original variance in the data?

3 PCs (PC3)


Q6. How many principal components (PCs) are required to describe at least 90% of the original variance in the data?

7 PCs (PC7)


Q7. What stands out to you about this plot? Is it easy or difficult to understand? Why?

This plot is difficult to understand, as the points are too close together to see any meaningful correlation.

```
biplot(wisc.pr)
```




```
# Plot PC2 vs PC1
plot(wisc.pr$x[,1:2], col=diagnosis, xlab="PC1", ylab="PC2")
```



> Q8. Generate a similar plot for principal components 1 and 3. What do you notice about these plots?

```
plot(wisc.pr$x[,c(1, 3)], col=diagnosis, xlab="PC1", ylab="PC3")
```

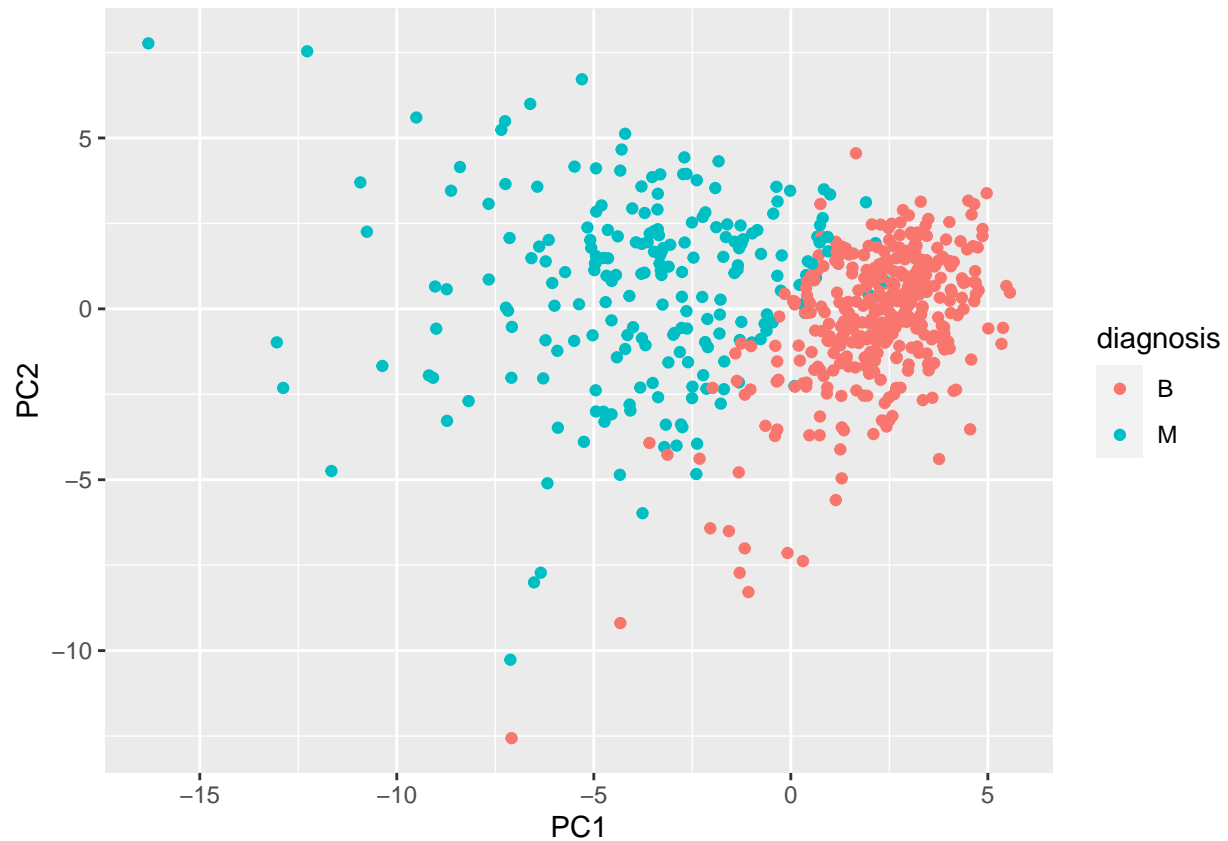


These plots display a distinction between two subgroups. There is a relatively clean separation between Benign and Malignant cell groups in the PC2 vs PC1 plot, while there is a less clean but still visible distinction in the PC3 vs PC1 plot. This is due to the fact that PC2 explains more variance in the data than PC3.

We can also create a plot using ggplot:

```
library(ggplot2)
wisc.df <- as.data.frame(wisc.pr$x)
ggplot(wisc.df, aes(wisc.pr$x[,1], wisc.pr$x[,2], col=diagnosis)) +
  geom_point() +
  labs(x="PC1", y="PC2")
```





Making scree plots to show how much variance is explained as the PC number increases:


First, store the variance in a variable.

```
pr.var <- wisc.pr$sdev^2
head(pr.var)
```

```
## [1] 13.281608  5.691355  2.817949  1.980640  1.648731  1.207357
```

Next, calculate the variance described by each PC vs that of all PCs.

```
pve <- (wisc.pr$sdev^2)/sum(wisc.pr$sdev^2)
plot(pve,
     xlab="Principal Component", ylab="Proportion of Variance Explained", type="o")
```

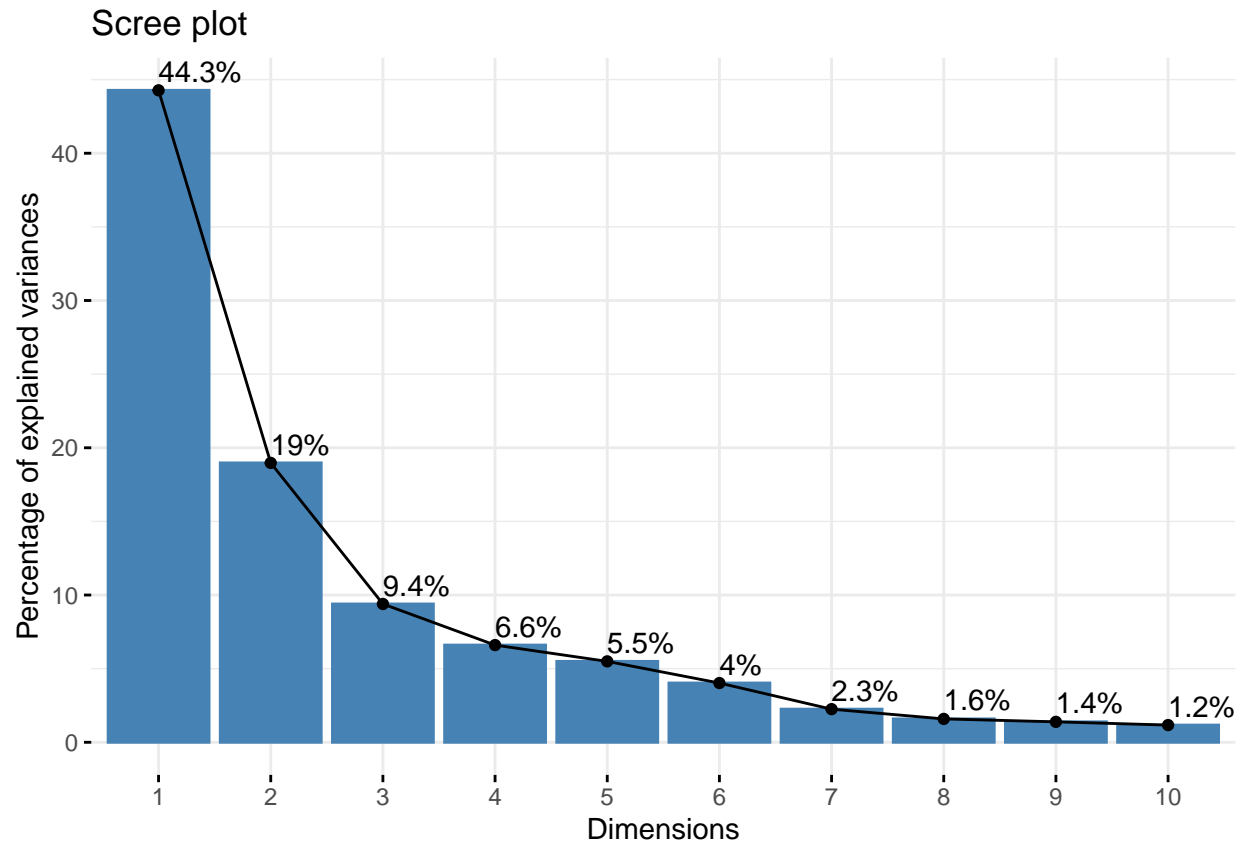


We can also use a barplot to represent this variance; this time we'll use the CRAN package `factoextra`.

```
library(factoextra)
```

```
## Welcome! Want to learn more? See two factoextra-related books at https://goo.gl/ve3WBa
```

```
fviz_eig(wisc.pr, addlabels=TRUE)
```



## Examine the PC loadings

How much do the original variables contribute to the new PCs that we have calculated? To get at this, we can look at the `$rotation` component of the returned PCA object.

```
head(wisc.pr$rotation[,1:3])
```

```
##           PC1          PC2          PC3
## radius_mean -0.2189024  0.23385713 -0.008531243
## texture_mean -0.1037246  0.05970609  0.064549903
## perimeter_mean -0.2275373  0.21518136 -0.009314220
## area_mean     -0.2209950  0.23107671  0.028699526
## smoothness_mean -0.1425897 -0.18611302 -0.104291904
## compactness_mean -0.2392854 -0.15189161 -0.074091571
```

Focus in on PC1

```
head(wisc.pr$rotation[,1])
```

```
##      radius_mean      texture_mean      perimeter_mean      area_mean
##      -0.2189024      -0.1037246      -0.2275373      -0.2209950
## smoothness_mean compactness_mean
##      -0.1425897      -0.2392854
```

Q9. For the first principal component, what is the component of the loading vector (i.e. `wisc.pr$rotation[,1]`) for the feature `concave.points_mean`?


```
wisc.pr$rotation["concave.points_mean",1]
```

```
## [1] -0.2608538
```

There is a complicated mix of variables that go together to make up PC1 (i.e. there are many of the original variables that together contribute highly to PC1).

```
loadings <- as.data.frame(wisc.pr$rotation)
```

```
ggplot(loadings) +  
  aes(PC1, rownames(loadings)) +  
  geom_col()
```



Q10. What is the minimum number of principal components required to explain 80% of the variance of the data?

```
summary(wisc.pr)
```

```
## Importance of components:
```

```
##          PC1      PC2      PC3      PC4      PC5      PC6      PC7
```

```
## Standard deviation      3.6444 2.3857 1.67867 1.40735 1.28403 1.09880 0.82172
## Proportion of Variance 0.4427 0.1897 0.09393 0.06602 0.05496 0.04025 0.02251
## Cumulative Proportion 0.4427 0.6324 0.72636 0.79239 0.84734 0.88759 0.91010
##          PC8      PC9      PC10      PC11      PC12      PC13      PC14
## Standard deviation      0.69037 0.6457 0.59219 0.5421 0.51104 0.49128 0.39624
## Proportion of Variance 0.01589 0.0139 0.01169 0.0098 0.00871 0.00805 0.00523
## Cumulative Proportion 0.92598 0.9399 0.95157 0.9614 0.97007 0.97812 0.98335
##          PC15      PC16      PC17      PC18      PC19      PC20      PC21
## Standard deviation      0.30681 0.28260 0.24372 0.22939 0.22244 0.17652 0.1731
## Proportion of Variance 0.00314 0.00266 0.00198 0.00175 0.00165 0.00104 0.0010
## Cumulative Proportion 0.98649 0.98915 0.99113 0.99288 0.99453 0.99557 0.9966
##          PC22      PC23      PC24      PC25      PC26      PC27      PC28
## Standard deviation      0.16565 0.15602 0.1344 0.12442 0.09043 0.08307 0.03987
## Proportion of Variance 0.00091 0.00081 0.0006 0.00052 0.00027 0.00023 0.00005
## Cumulative Proportion 0.99749 0.99830 0.9989 0.99942 0.99969 0.99992 0.99997
##          PC29      PC30
## Standard deviation      0.02736 0.01153
## Proportion of Variance 0.00002 0.00000
## Cumulative Proportion 1.00000 1.00000
```

5 PCs (PC5)

## Hierarchical Clustering

Scale the data

```
data.scaled <- scale(wisc.data)
```


Calculate Euclidean distances between all pairs of observations

```
data.dist <- dist(data.scaled)
```

Q11. Using the `plot()` and `abline()` functions, what is the height at which the clustering model has 4 clusters?

```
wisc.hclust <- hclust(data.dist, )
plot(wisc.hclust)
abline(h=19, col="red", lty=2)
```

## Cluster Dendrogram



```
data.dist
hclust(*, "complete")
```

A height of 19 yields 4 clusters for the model.

Cut this tree to yield cluster membership vector using `cutree()` function.

```
wisc.hclust.clusters <- cutree(wisc.hclust, h=19)
table(wisc.hclust.clusters)
```

```
## wisc.hclust.clusters
##  1  2  3  4
## 177  7 383  2
```

```
table(wisc.hclust.clusters, diagnosis)
```

```
##              diagnosis
## wisc.hclust.clusters  B  M
##              1  12 165
##              2   2   5
##              3 343  40
##              4   0   2
```

Q12. Can you find a better cluster vs diagnoses match by cutting into a different number of clusters between 2 and 10?

```
wisc.hclust.clusters.4 <- cutree(wisc.hclust, k=4)
table(wisc.hclust.clusters.4, diagnosis)
```

```
##              diagnosis
## wisc.hclust.clusters.4  B  M
##              1  12 165
##              2   2   5
##              3 343  40
##              4   0   2
```

The best cluster vs diagnosis match is at k=4 clusters, as it is the lowest number of groupings necessary to separate the benign and malignant groups accurately enough, before the increase in accuracy tapers off as the value of k increases.

Q13. Which method gives your favorite results for the same data.dist dataset? Explain your reasoning.

```
methods <- function(x) {
  table(cutree(hclust(data.dist, method=x), k=4), diagnosis)
}

methods("single")
```

```
##      diagnosis
##        B  M
##  1 356 209
##  2   1   0
##  3   0   2
##  4   0   1
```

```
methods("complete")
```

```
##      diagnosis
##        B  M
##  1  12 165
##  2   2   5
##  3 343  40
##  4   0   2
```

```
methods("average")
```

```
##      diagnosis
##        B  M
##  1 355 209
##  2   2   0
##  3   0   1
##  4   0   2
```

```
methods("ward.D2")
```

```
##      diagnosis
##        B  M
##  1   0 115
##  2   6  48
##  3 337  48
##  4  14   1
```

method="complete" and method="ward.D2" give my favorite results. "ward.D2" generally separates B and M into two groups, with some overlap and increased clustering in other groups; "complete" in this case is the best method, since it keeps clustering mostly within two groups, 1 and 3, and separates B and M about as well as "ward.D2". In contrast, "single" and "average" both leave nearly all the data points in one group, failing to separate the diagnoses at all.

## K-means clustering

```
wisc.km <- kmeans(data.scaled, centers=2, nstart=50)
table(wisc.km$cluster, diagnosis)
```

```
##      diagnosis
##           B   M
##    1 343   37
##    2  14  175
```

Q14. How well does k-means separate the two diagnoses? How does it compare to your hclust results?

K-means does a good job of separating the two diagnoses, producing very similar results to the hclust results. There is a consistently high amount of false positives and a lower amount of false negatives in both models.

Comparing clusters using k-means model with clusters using hierarchical clustering model:

```
table(wisc.hclust.clusters, diagnosis)
```


```
##              diagnosis
## wisc.hclust.clusters  B   M
##           1  12 165
##           2   2   5
##           3 343  40
##           4   0   2
```

## Combine methods: PCA and HCLUST

My PCA results were interesting as they showed a separation of M and B samples along PC1.

```
plot(wisc.pr$x[,1:2], col=diagnosis)
```





I want to cluster my PCA results - that is use `wisc.pr$x` as input to `hclust`.


Try clustering in 3 PCs, that is PC1, PC2, PC3 as input

```
d <- dist(wisc.pr$x[,1:3])  
wisc.pr.hclust <- hclust(d, method="ward.D2")
```

Tree result figure:

```
plot(wisc.pr.hclust)
```

## Cluster Dendrogram




d  
hclust (\*, "ward.D2")

Let's cut this tree into two groups/clusters

```
wisc.hclust.clusters <- cutree(wisc.pr.hclust, k=2)
table(wisc.hclust.clusters)
```

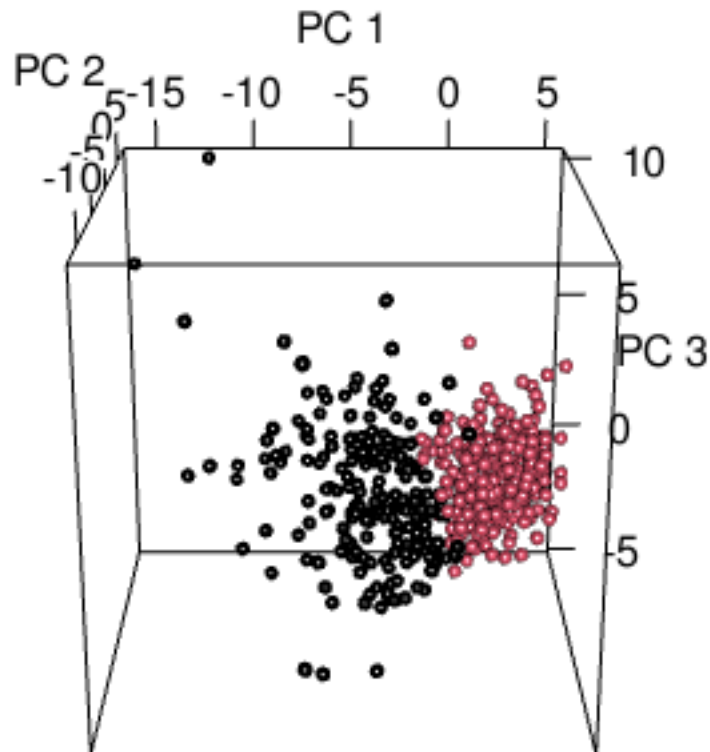
```
## wisc.hclust.clusters
##    1    2
## 203 366
```

```
plot(wisc.pr$x[,1:2], col=wisc.hclust.clusters)
```



```
library(rgl)
plot3d(wisc.pr$x[,1:3], xlab="PC 1", ylab="PC 2", zlab="PC 3", cex=1.5, size=1, type="s", col=wisc.hclu
rglwidget(width = 400, height = 400)
```

```
## Warning in snapshot3d(scene = x, width = width, height = height): webshot =
## TRUE requires the webshot2 package and Chrome browser; using rgl.snapshot()
## instead
```



How well do the two clusters separate the M and B diagnosis?

```
table(wisc.hclust.clusters, diagnosis)
```

```
##              diagnosis
## wisc.hclust.clusters  B  M
##              1  24 179
##              2 333  33
```

Calculate the accuracy of our results

```
(179+333)/nrow(wisc.data)
```

```
## [1] 0.8998243
```

Nearly 90% of our results are accurate. However, a 10% false positive could potentially be dangerous to patients.

Q15. How well does the newly created model with four clusters separate out the two diagnoses?

```
wisc.pr.hclust.clusters.4 <- cutree(wisc.pr.hclust, k=4)
table(wisc.pr.hclust.clusters.4, diagnosis)
```

```
##              diagnosis
## wisc.pr.hclust.clusters.4  B  M
##              1   0 111
##              2  24  68
##              3 184  32
##              4 149   1
```

With four clusters, this new model produces a group with too much of each diagnosis to be as useful as with 2 clusters.

Q16. How well do the k-means and hierarchical clustering models you created in previous sections (i.e. before PCA) do in terms of separating the diagnoses? Again, use the `table()` function to compare the output of each model (`wisc.km$cluster` and `wisc.hclust.clusters`) with the vector containing the actual diagnoses.

```
table(wisc.km$cluster, diagnosis)
```

```
##      diagnosis
##           B   M
##  1 343  37
##  2  14 175
```

```
wisc.pr.hclust.clusters.2 <- cutree(wisc.pr.hclust, k=2)
table(wisc.pr.hclust.clusters.2, diagnosis)
```

```
##              diagnosis
## wisc.pr.hclust.clusters.2  B   M
##              1  24 179
##              2 333  33
```

The k-means and hierarchical clustering combination separated the diagnoses with similar accuracy as the PCA clustering model.

## Sensitivity and specificity

Q17. Which of your analysis procedures resulted in a clustering model with the best specificity? How about sensitivity?

```
# function for sensitivity and specificity calculations
sens.spec <- function(tp, fn, tn, fp) {
  sens <- tp/(tp+fn)
  spec <- tn/(tn+fp)
  results <- c(sens, spec)
  names(results) <- c("Sensitivity", "Specificity")
  results
}
```

```

# results for each method
## table(wisc.hclust.clusters, diagnosis)
hclust <- sens.spec(165, 12, 343, 40)

## table(wisc.km$cluster, diagnosis)
kmeans <- sens.spec(175, 14, 343, 37)

## table(wisc.pr.hclust.clusters.2, diagnosis)
pca <- sens.spec(179, 24, 333, 33)

# data frame of all method results
sens.spec.df <- data.frame(hclust, kmeans, pca)
sens.spec.df

```

```

##           hclust    kmeans      pca
## Sensitivity 0.9322034 0.9259259 0.8817734
## Specificity 0.8955614 0.9026316 0.9098361

```

```

# calculate colnames of max sensitivity and max specificity
sens.spec.df$max <- colnames(sens.spec.df)[apply(sens.spec.df, 1, which.max)]
best <- sens.spec.df$max
names(best) <- c("Best Sensitivity", "Best Specificity")
best

```

```

## Best Sensitivity Best Specificity
##           "hclust"           "pca"

```

## Prediction

```

url <- "https://tinyurl.com/new-samples-CSV"
new <- read.csv(url)
npc <- predict(wisc.pr, newdata=new)
npc


```

```

##           PC1      PC2      PC3      PC4      PC5      PC6      PC7
## [1,]  2.576616 -3.135913  1.3990492 -0.7631950  2.781648 -0.8150185 -0.3959098
## [2,] -4.754928 -3.009033 -0.1660946 -0.6052952 -1.140698 -1.2189945  0.8193031
##           PC8      PC9      PC10     PC11     PC12     PC13     PC14
## [1,] -0.2307350 0.1029569 -0.9272861 0.3411457  0.375921 0.1610764 1.187882
## [2,] -0.3307423 0.5281896 -0.4855301 0.7173233 -1.185917 0.5893856 0.303029
##           PC15     PC16     PC17     PC18     PC19     PC20
## [1,] 0.3216974 -0.1743616 -0.07875393 -0.11207028 -0.08802955 -0.2495216
## [2,] 0.1299153  0.1448061 -0.40509706  0.06565549  0.25591230 -0.4289500
##           PC21     PC22     PC23     PC24     PC25     PC26
## [1,] 0.1228233 0.09358453 0.08347651 0.1223396 0.02124121 0.078884581
## [2,] -0.1224776 0.01732146 0.06316631 -0.2338618 -0.20755948 -0.009833238
##           PC27     PC28     PC29     PC30
## [1,] 0.220199544 -0.02946023 -0.015620933 0.005269029
## [2,] -0.001134152 0.09638361 0.002795349 -0.019015820

```

```
plot(wisc.pr$x[,1:2], col=wisc.hclust.clusters)
points(npc[,1], npc[,2], col="blue", pch=16, cex=3)
text(npc[,1], npc[,2], c(1,2), col="white")
```



Q18. Which of these new patients should we prioritize for follow up based on your results?

We should prioritize patient 2, as their cell data seem to correspond to the PCA group denoting a malignant diagnosis.